

Aya Alsadi

101115964

STAT 4601

Data Mining

Final Project

Due: Dec-28-2020

## PART I. DATA Visualization

### 1. The data set:

The file audit.csv contains characteristics of 2000 individual tax returns. The dependent variables are (TARGET\_Adjusted) and (RISK\_Adjustment). The data set contains a total of 12 variables including: ID (Unique identifier for each person), Age (Age of person), Employment (Type of employment), Education (Highest level of education), Marital (Current marital status), Occupation (Type of occupation), Income (Amount of income declared), Gender (Gender of person), Deductions (Total amount of expenses that a person claims in their financial statement), Hours (Average hours worked on a weekly basis), Risk\_Adjustment (monetary amount of any adjustment to the person's financial status as a result of a productive audit. This is a measure of the size of the risk associated with the person), TARGET\_Adjusted (binary target variable for classification modeling (0/1), indicating non-productive and productive audits, respectively. Productive audits are those that result in an adjustment being made to a client's financial statement.)

The excel file was loaded from RStudio. the original data is shown below:

```
> library(readxl)
> AuditData <- read_excel("C:/Users/olivierM/Desktop/AuditData.xlsx")
> View (AuditData)
```

	ID	Age	Employment	Education	Marital	Occupation	Income	Gen	Hrs	Ded	Risk	Target
1	1004641	38	Private	College	Unmarried	Service	81838.00	M				
2	1010229	35	Private	Associate	Absent	Transport	72099.00	M				
3	1024587	32	Private	HSgrad	Divorced	Clerical	154676.74	M				
4	1038288	45	Private	Bachelor	Married	Repair	27743.82	M				
5	1044221	60	Private	College	Married	Executive	7568.23	M				
6	1047095	74	Private	HSgrad	Married	Service	33144.40	M				
7	1047698	43	Private	Bachelor	Married	Executive	43391.17	M				
8	1053888	35	Private	Yr12	Married	Machinist	59906.65	M				
9	1061323	25	Private	Associate	Divorced	Clerical	126888.91	F				
10	1062363	22	Private	HSgrad	Absent	Sales	52466.49	F				
11	1068642	48	Private	College	Divorced	Service	291416.11	F				

The data can be treated as a matrix. Since this is a large set (n=2000 units), it is not convenient to display all in the R window. Parts of the matrix are shown below:

```
> mdata=AuditData;
> mdata[1:5,1:7]
```

```

# A tibble: 5 x 7
   ID    Age Employment Education Marital Occupation Income
   <dbl> <dbl> <chr>      <chr>    <chr>    <chr>    <dbl>
1 1004641  38 Private    College  Unmarried Service  81838
2 1010229  35 Private    Associate Absent   Transport 72099
3 1024587  32 Private    HSgrad    Divorced Clerical 154677.
4 1038288  45 Private    Bachelor Married  Repair   27744.
5 1044221  60 Private    College  Married  Executive 7568.

> mdata[10:15,2:8]
# A tibble: 6 x 7
   Age Employment Education Marital Occupation Income Gender
   <dbl> <chr>      <chr>    <chr>    <chr>    <dbl> <chr>
1 22 Private    HSgrad    Absent   Sales     52466. Female
2 48 Private    College  Divorced Service  291416. Female
3 60 Private    Vocational Widowed Clerical 24155. Male
4 21 Private    College  Absent   Service  143255. Female
5 21 Private    College  Absent   Machinist 120555. Male
6 50 Private    Master   Married  Executive 34919. Male

```


## 2. A visual look into the data

The second command shows a matrix of scatters for five quantitative variables. Each scatter describes the degree of linear relationship between two variables.

```

> mData2 = cbind (Age, Income, Hours, RiskAdj, Deductions);
> pairs(mData2)


```





### 3. Conditional Plots


To investigate further some of the more complicated relationships we can look at conditional plotting. We present several plots corresponding to different ranges of the conditional variable. For instance, we plot “Age” against “Income”, given “Hours”. We also look at other conditional relationships’ as shown below.

```
> coplot(Age~Income|Hours)
> coplot(Age~Income|RiskAdj)
> coplot(Age~Income|Deductions)
> coplot(Hours~Income|Deductions)
> coplot(Hours~Income|RiskAdj)
```




In both the figure above, the lower six panels show the pairwise plots for “Age” against “Income” for different ranges of “Hours” as shown in the upper panel. The default function selects 6 different subsets of the third variable “Hours” with an equal number of cases in each. In addition, an overlap of 0.5 is allowed.





#### 4. Data Visualization in Ggobi

We continue our investigation of the data using the program Ggobi. Data Visualization in Ggobi can be done thru the use of a package that may be called from R or used alone. In this work, we used Ggobi alone independently of the R program. The excel data Audit.csv was first loaded from the console.




Ggobi is well suited for visualizing high dimensional data and dynamic graphics and for investigating hidden relationships.


## 1D graphics

In order to investigate the data, we will start with a “tour” in 1D. This represents the projection of the 12 variables/axes on a 1-dimensional display. We begin with 1D plots of the quantitative variables in the data set.

*Histograms and Bar charts:* These types of chart are well suited for visualising the distribution of qualitative data (bar chart) and quantitative data (histogram). We would like to know if our data is skewed to the left or right or relatively symmetric, and whether there are severe outliers that could be removed during the preprocessing (or cleaning) phase.

**Figure 1.** Histograms and Bar charts.





*Comments:*

The graphs reveal a few facts about the data. The variables “Age”, “Deduction”, “Risk adjustment”, and “Income” are skewed to the left. The variable “Hours” is relatively symmetric. As for the categorical items, “Gender” and “Marital”, there appear to be more male than female in the sample, and more married or **absent** individuals in the sample with respect to “Marital status”.

**Formatted:** Highlight


## 2D plots: Scatter matrix

This represents the projection of the 12 variables/axes on a 2-dimensional display. The procedure used in the grand tour projects the data by selecting two directions. This allows the observations of the data from all directions.


A common 2D plot is the Scatterplot Matrix. This plots the variables pairwise and allows to see if there are some types of relationships between variables by displaying pairwise plots.

High correlations (close to straight line scatter) indicate some linear relationship between pairs of variables but low correlations may indicate either no relationship or the presence of some nonlinear relationship in combination with other variables.

**Figure 2.** Scatterplot matrix for “Age” and “Income”.




**Figure 3.** Scatterplot matrix for “Age”, “Employment”, “Education” and “Id”.



#### High dimensional data visualization.

The Parallel Coordinates Plot is a method for investigating high dimensional data. Each data point has a value on each of the axes which are plotted vertically rather than at right angles to each other. We show in figure 4 below, two parallel coordinates plots for two subsets of the data set.

**Figure 4.** Parallel Coordinates Plot



**Conclusion:**

This section was meant to get a visual feeling of the data set “Audit.csv”. we presented 1D plots for both numerical and categorical variables. We also investigated relationships between in the data with 2D graphs. The use of Ggobi was useful in getting in-depth informations about the distributions of the variables and their relationships. Although this set contains few variables, dimension reduction is a key step in the multivariable analysis. In the next section, we shall select a subset of the variables in the data set and apply a dimension reduction technique known as PCA (principal component analysis).

## PART II. Dimension reduction

The aim of dimension reduction is to compress the number of variables in the original data into a small set so as to ease the subsequent data analyses (such as regression, correlation analyses, hypotheses testing, etc.) that might be performed on the data. This step, although crucial in general, does appear non pertinent in our study of the data “Audit.csv”. Indeed, there are 6 quantitative variables and only 5 pertinent for such analysis. Thus, this section purpose is to demonstrate the application of a classic dimension reduction algorithm, namely PCA (Principal Component Analysis). Although purely exploratory, PCA results can be used as a basis for investigating other data. Thus, one limitation of this section is the inability to generalize the result of this study.

### 1. Methodology:

For the purpose of this analysis, we selected a subset of 5 variables: “Age”, “Income”, “Hours”, “Deductions” and “Risk adjustment”. These variables are all measured on a numerical scale.

*Preprocessing:* The missing values were first removed and the data scaled before the cluster analysis.

In R, the first three cases of the set are shown below:

```
> Deductions=AuditData$Deductions;
> Data1=cbind(Age, Income, Hours, RiskAdj, Deductions);
> head(Data1)

> df=Data1;
> df=na.omit(df);
> df=scale(df);
> head(df, n = 3)
      Age      Income      Hours      RiskAdj Deductions
[1,] -0.04578664 -0.04094215  2.626809833 -0.2422672 -0.1983193
[2,] -0.26662253 -0.18082681 -0.828923452 -0.2422672 -0.1983193
[3,] -0.48745842  1.00526611 -0.006129813 -0.2422672 -0.1983193

> dim(df)
[1] 2000     5
```

### 2. Running the PCA algorithm

With a total of p=5 variables, exactly 5 principals' components were extracted, along with their standard deviation and proportion of variance.

```
> data.pca=prcomp(df, center = TRUE, scale. = TRUE);
> summary(data.pca)

Importance of components:
                    PC1       PC2       PC3       PC4       PC5 
Standard deviation   1.1964   1.0103   0.9632   0.9581   0.8380 
Proportion of Variance 0.2863   0.2041   0.1855   0.1836   0.1404 
Cumulative Proportion  0.2863   0.4904   0.6760   0.8596   1.0000
```

The variance of each component is measured by the eigenvalue for that component. The higher the eigenvalue, the better the component captures the total information in the original data set.

```
> eig.val=get_eigenvalue(data.pca)
> eig.val
  eigenvalue variance.percent cumulative.variance.percent
Dim.1  1.4313662      28.62732          28.62732
Dim.2  1.0206774      20.41355          49.04987
Dim.3  0.9277295      18.55459          67.59546
Dim.4  0.9180146      18.36029          85.95575
Dim.5  0.7022124      14.04425         100.00000
```

*Comments:*


The first 4 components amount to a total of 85.96% of the total variance. This means that about 85.96% of the variation or information in the original data set is captured altogether by the first 4 (best) components. The first three component capture about 68% of the total variance. That is, by reducing the dimension from 5 to 3, we capture 67.6% of the total information. A look at the table above shows that the last component contributes very little in term of percentage to the total variance (about 14%).

How many components must be retained? The use of the scree plot can help answer this question.

### **3. Scree plot, for deciding on the number of components**


The scree plot graphs the variance on the y-axis and the number of the component on the x-axis. It shows the variation of the variance captures as we increase the number of components. In general, less variance is captured as the number of components is included in the model. This allows for dimension reduction (we can suppress those components with small contribution in terms of variance).

```
> screeplot(data.pca,main = "Screeplot of the first PCs")
> abline(h = 1, col="red", lty=5)
> legend("topright", legend=c("Eigenvalue = 1"),col=c("red"), lty=5, cex=0.6)
> fviz_eig(data.pca)
```



numbered and there is not label on the horizontal axis for this one.


Graphs should be



The cumulative variance plot is also helpful in determining the appropriate number of principal components.

The graph below shows the cumulative plot for

```
> cumpro=cumsum (data.pca$sdev^2 / sum(data.pca$sdev^2))
> plot(cumpro[0:15], xlab = "PC #", ylab = "Amount of explained variance", main = "Cumulative variance plot")
> abline(v = 2, col="blue", lty=5)
> abline(h = 0.5, col="blue", lty=5)
> legend ("topleft", legend=c ("Cut-off @ PC1"), col=c("blue"), lty=5, cex=0.6)
```




*Comments: The scree plots suggest that 3 or 4 components could be optimal. However, the original number of variables ( $p=5$ ) is too small to make this finding pertinent.*

#### 4. Plot of the first two principal components

The plot below shows the representation of the data points in the first 2 components plane.

```
> plot(data.pca$x[,1], data.pca$x[,2], xlab="PC1", ylab = "PC2", main = "PC1  
vs. PC2 - plot")
```



The coordinates (or scores) of the participants on the components are:

```
> data.pca$x[1:5,1:5]
      PC1     PC2     PC3     PC4     PC5
[1,] 0.9836440 1.73169248 -0.5056904 1.13916939 -1.2174492
[2,] -0.5281612 -0.17170175 0.3580401 -0.27738971 0.6126585
[3,] -0.9845120  0.06628008 -0.3592115  0.13929232 -0.4734396
[4,]  1.4410183  0.78664207 -0.3435072 -0.01416525 -0.1533557
[5,]  2.0027353 -0.34889351 -0.1786144 -1.40343912  0.1958866
```

The correlation between the variables and the components are:

```
> data.pca
Standard deviations (1, ... , p=5):
[1] 1.1963972 1.0102858 0.9631871 0.9581308 0.8379811

Rotation (n x k) = (5 x 5):
          PC1        PC2        PC3        PC4        PC5
Age      0.5095489 -0.2598393  0.430095139 -0.45792095 -0.5274182970
Income    -0.5872600 -0.2483125 -0.341786131 -0.03972703 -0.6892536830
Hours     0.4286868  0.5774494 -0.263591025  0.44078682 -0.4679813792
RiskAdj   0.3865196 -0.2244748 -0.792917190 -0.37913576  0.1665895460
Deductions 0.2496316 -0.6978450 -0.001585471  0.67133689  0.0008076244
```

High correlations (close to 1 or -1) indicate a strong relationship between the variable and the component. That is, the component captures the variability in that variable, for instance, the matrix above shows that Risk adjustment has a large loading on PC3 (-0.79) while Income has a very low relationship with PC4 (-0.0397).

### *5. Quality of representation*

The representation of a variable on a component can be measured with the cosine squared index and its contribution to the component.

```
> res.var=get_pca_var(data.pca);
> x1=res.var$cos2[,1:5];
> x2=res.var$contrib[,1:5];

> x1
           Dim.1       Dim.2       Dim.3       Dim.4       Dim.5
Age      0.37163998 0.06891253 1.716131e-01 0.192499941 1.953345e-01
Income    0.49364134 0.06293406 1.083753e-01 0.001448844 3.336005e-01
Hours     0.26304552 0.34034269 6.445886e-02 0.178363820 1.537891e-01
RiskAdj   0.21384238 0.05143085 5.832799e-01 0.131959017 1.948785e-02
Deductions 0.08919696 0.49705732 2.332051e-06 0.413742935 4.580230e-07

> x2
           Dim.1       Dim.2       Dim.3       Dim.4       Dim.5
Age      25.964005 6.751647 1.849818e+01 20.9691599 2.781701e+01
Income    34.487425 6.165911 1.168178e+01 0.1578237 4.750706e+01
Hours     18.377234 33.344784 6.948023e+00 19.4293019 2.190066e+01
RiskAdj   14.939740 5.038893 6.287177e+01 14.3743927 2.775208e+00
Deductions 6.231596 48.698766 2.513719e-04 45.0693218 6.522571e-05
```

The large numbers mean the item is well represented on the component.

*Interpretation of the best 2 components:*

We first rank the quality of representation of each of the 42 items of the first dimension (from large to small).

```
> x1=res.var$cos2;
> sort(x1, decreasing=TRUE)
```

5.832799e-01 4.970573e-01 4.936413e-01 4.137429e-01 3.716400e-01 3.403427e-01  
2.630455e-01 2.138424e-01 1.953345e-01 1.924999e-01 1.783638e-01 1.716131e-01  
1.319590e-01 1.083753e-01 8.919696e-02 6.891253e-02 6.445886e-02 6.293406e-02  
1.948785e-02 1.448844e-03 2.332051e-06 4.580230e-07

Dimension#1: The first dimension is strongly correlated with the following items: (*Age, Income*), so the first dimension can be labelled SOCIO-ECONOMIC STATUS.

Dimension#2: The second dimension is strongly correlated with the following items (*Hours, Deductions*). “Deductions” is the total amount of expenses that a person claims in their financial statement, while “Hours” denotes the “Average hours worked on a weekly basis”. So, the second dimension can be labelled FINANCIAL STATUS.

### **Conclusion:**

We performed data reduction on a group of 5 numerical variables from the data set Audit.csv. Principal Component Analysis was used to extract 5 components, from which about 3 or 4 seem to be adequate to represent the original data. With only 3 components, we captured 67.6% of the total variance. With 4 components, about 85% of the total variance was captured. The main limitation of this analysis lies in the fact that the original data set is not well suited for this type of analysis. Dimension reduction is usually applied on large data sets with hundreds of dimensions or variables.

## **PART III. Data reduction. Clustering.**

### ***I. The data set and reason for Clustering:***

The Audit data file contains characteristics of 2000 individual tax returns. The variables include Age (Age of person), Education (Highest level of education), Outcome (Amount of income declared), Deductions (Total amount of expenses that a person claims in their financial statement), Hours (Average hours worked on a weekly basis), Risk Adjustment (monetary amount of any adjustment to the person’s financial status as a result of a productive audit).

Our data set contains an overwhelming amount of data in terms of number of cases ( $n=2000$ ). We want to consider reducing the sample size  $n$  (number of cases). One methodology is to cluster together similar objects or cases. The goal is to group together sample units that look alike in such a way as to increase the degree of homogeneity within each cluster and heterogeneity between the clusters.

## **2. Kmeans as a clustering algorithm**

Clustering is an unsupervised learning, because we do not know which class or cluster an observation belongs to. Various algorithms exist for cluster analysis. We shall use the commonly used Kmeans algorithm,

K-means clustering is an unsupervised machine learning algorithm for partitioning a given data set into a set of k groups or clusters, where k represents the number of groups pre-specified by the analyst. We applied the kmeans algorithm to the Audit data.

## **3. Methodology:**

We selected a subset of features, because the Kmeans procedure works well on numerical attributes. The selected variables were combined in one data set, namely “Data1” which includes “Age”, “Income”, “Hours”, and “Risk adjustment”. The classification of observations into groups requires some methods for computing the distance or the (dis)similarity between each pair of observations. We use the Euclidean distance for simplicity.

### *Preprocessing:*

The missing values were first removed and the data scaled before the cluster analysis.


```
> df>Data1;
> df=na.omit(df);
> df=scale(df);
> head(df, n = 3)
      Age      Income      Hours      RiskAdj
[1,] -0.04578664 -0.04094215  2.626809833 -0.2422672
[2,] -0.26662253 -0.18082681 -0.828923452 -0.2422672
[3,] -0.48745842  1.00526611 -0.006129813 -0.2422672
```

## **4. Selecting the optimal number of clusters**



*The Kmeans procedures takes an integer K as input, for the number of categories in the data.*  
We applied 3 different methods for selecting the right number K of clusters in our data set. The “Elbow method”, the “Silhouette method”, and the “Gap statistic method”. Unfortunately, the Gap method did not converge for our data set.

The other two methods produced different results, Elbow method (k = 4), Silhouette method (k=7). We decided to be conservative and chose a value close to the average of the two methods recommendations. We chose k = 5 for our analysis.

**Figure 1.** Optimal number of clusters. “Elbow method”



**Figure 2.** Optimal number of clusters. “Silhouette and Gap methods”



5. Running Kmeans, for  $k = 5$

The k-means algorithm tries to cluster objects that are ‘close’ together, using a distance measure. The following will consider Euclidean distance. The algorithm works as follow:

1. Select a number of points K to be the centers (centroids)
2. Start with a random initialization for the centers.
3. Find the distance from each point to every center; assign the point to closest center;
4. For each set of points assigned to a center, find the centroid of the cluster;
5. Take that value as the new center;
6. Repeat the process until the centers do not seem to move.

*The results in R are as follow:*

```
> set.seed (123)
> km.res=kmeans (df, 5, nstart = 25)
> print(km.res)

K-means clustering with 5 clusters of sizes 367, 355, 712, 12, 554
```

Cluster means:

	Age	Income	Hours	RiskAdj
1	0.05650791	-0.4765976	1.3495740	0.125844874
2	-0.48061562	1.7569470	-0.4729203	-0.170213252
3	-0.69413021	-0.2567961	-0.2556258	-0.180245915
4	0.69646732	-0.3109788	0.2818480	11.531038989
5	1.14755098	-0.4733471	-0.2685624	0.007587111

Clustering vector:

```
[1] 1 3 2 1 5 5 1 3 3 3 2 5 2 3 5 3 3 3 5 1 2 3 3 3 2 3 3 1 5 3 3 3 2 5
[35] 3 2 2 2 5 3 5 5 3 1 3 5 5 3 5 2 1 3 3 5 3 3 3 3 1 1 5 3 5 5 1 3 5 1
[69] 5 3 3 3 5 5 3 5 1 3 2 1 1 1 3 1 2 5 3 1 3 1 3 5 5 5 2 5 1 3 3 5 2 1
[103] 5 3 1 5 5 3 1 5 5 5 5 3 2 1 5 5 3 1 3 2 3 5 1 2 3 2 3 3 2 3 5 1 2
[137] 1 3 5 3 5 3 3 5 1 3 3 2 3 1 3 1 5 3 2 5 3 3 2 3 5 5 1 3 5 5 3 1 5 5
[171] 5 5 5 3 3 5 3 1 1 1 3 1 3 3 3 1 3 5 3 3 5 5 1 5 3 5 3 2 3 2 3 3 3 3
[205] 5 3 3 3 1 3 1 5 3 3 2 5 3 2 5 3 1 5 5 3 3 3 3 3 5 2 1 5 5 5 5 3 1 1
[239] 2 3 2 1 3 3 5 5 3 3 3 1 5 3 1 5 1 2 3 2 3 5 2 1 3 5 1 5 5 3 5 2 3 5
[273] 2 5 3 5 3 2 3 1 5 5 5 1 5 1 2 5 2 3 5 3 3 2 1 5 3 3 3 3 1 3 2 3 1
```

```
within cluster sum of squares by cluster: [1] 630.25876 760.94887 712.64083 3
0.26892 878.73915(between_SS / total_SS = 62.3 %)
```

We aggregated the data set with the cluster number and mean, so that each cluster is shown with its mean for all four covariates.:

```
> aggregate (Data1, by=list(cluster=km.res$cluster), mean)
```

cluster	Age	Income	Hours	RiskAdj
1	39.38965	51506.95	56.47684	3070.7439
2	32.09296	207010.00	34.32676	601.0648
3	29.19242	66809.89	36.96770	517.3736
4	48.08333	63037.60	43.50000	98211.4167
5	54.21119	51733.26	36.81047	2084.2527

The table below shows the first six participants of the study with their scores on each variable (unscaled) and the cluster to which they belong.

```
> dd = cbind (Data1, cluster = km.res$cluster);  
> head(dd)
```

	Age	Income	Hours	RiskAdj	cluster
[1,]	38	81838.00	72	0	1
[2,]	35	72099.00	30	0	3
[3,]	32	154676.74	40	0	2
[4,]	45	27743.82	55	7298	1
[5,]	60	7568.23	40	15024	5
[6,]	74	33144.40	30	0	5

The next table shows the first six participants of the study with their scores on each variable (scale) and the cluster to which they belong.

```
> dd = cbind (df, cluster = km.res$cluster);  
> head(dd)
```

	Age	Income	Hours	RiskAdj	cluster
[1,]	-0.04578664	-0.04094215	2.626809833	-0.2422672	1
[2,]	-0.26662253	-0.18082681	-0.828923452	-0.2422672	3
[3,]	-0.48745842	1.00526611	-0.006129813	-0.2422672	2
[4,]	0.46949710	-0.81791576	1.228060646	0.6325964	1
[5,]	1.57367655	-1.10770480	-0.006129813	1.5587673	5
[6,]	2.60424404	-0.74034534	-0.828923452	-0.2422672	5


```
> head (km.res$cluster, 5)  
[1] 1 3 2 1 5
```

## 6. Graph of the clusters

We can get a visual look at the spatial configuration of the k=5 groupings. If we print the results, we see that our groupings resulted in 5 cluster sizes which are respectively: 367, 355, 712, 12, 55. We see the cluster centers (means) for the groups across the four variables “Age”, “Income”, “Hours”, and “Risk adjustment”. We also get the cluster assignment for each observation represented by its ID, as shown below.

```
> fviz_cluster (km.res, df)
```

**Figure 3.** Cluster plot on first two dimensions.




**Comments:** We have 5 relatively distinct clusters. The total within-cluster sum of square is a measure of the compactness (i.e., goodness) of the clustering (should be as small as possible). In our case, we have the following goodness fit indices (within sum = 3012.86).

```
> km.res$tot.withinss
[1] 3012.857
> km.res$betweenss
[1] 4983.143
```

## 7. Varying the number of clusters k

Because the number of clusters (k) must be set before we start the algorithm, it is often advantageous to use several different values of k and examine the differences in the results. We executed the same process for 2, 3, 4, and 5 clusters, and the results are shown in the figure below. This graph allows a comparison of different clustering models and appreciation of the improvement in the data reduction, as we varied the number of clusters.

```
> k3=kmeans(df, centers = 3, nstart = 25);
> k4=kmeans(df, centers = 4, nstart = 25);
> k5=kmeans(df, centers = 5, nstart = 25);
> k2=kmeans(df, centers = 2, nstart = 25);
> p1=fviz_cluster(k2, geom = "point", data = df) + ggtitle("k = 2")
> p2=fviz_cluster(k3, geom = "point", data = df) + ggtitle("k = 3")
> p3=fviz_cluster(k4, geom = "point", data = df) + ggtitle("k = 4")
> p4=fviz_cluster(k5, geom = "point", data = df) + ggtitle("k = 5")
> grid.arrange(p1, p2, p3, p4, nrow = 2)
```



*Comments:* Although k=5 clusters seem to be optimal, the data reduction is relatively good for even small clusters such as k=2 or k=3. We see that the goodness fit increases but very slowly as

the number of clusters increases. Indeed, the total within-cluster sum of square measures for each model is:

```
> wws = c(k2$tot.withinss, k3$tot.withinss, k4$tot.withinss, k5$tot.withinss)
> wws
[1] 6153.580 5171.976 3616.830 3012.857
```

#### **Conclusion:**

The K-means clustering procedure is simple and fast. Furthermore, it can efficiently deal with very large data sets. However, there are some weaknesses of the k-means approach.

The limitations include non reproducibility and the impossibility of solutions. Different runs on the same data set can lead to situations where some points can be assigned to different centers for different initializations. Thus, the algorithm can end up with different clusters after a new initialization. Moreover, if the initial center is moved to an empty region of the plane, the algorithm will fail to find the clusters. Although random initialization can resolve this issue. Another potential disadvantage of K-means clustering is that it requires to pre-specify the number of clusters. The method is also sensitive to outliers.

Finally, there is no definitive answer to the question of the optimal number of clusters. The number k is subjective and depends on the method used for measuring similarities and the parameters used for partitioning. Another solution consists of inspecting the dendrogram produced using hierarchical clustering to see if it suggests a particular number of clusters. However, this approached led to uninterpretable results for our data set.

In this section, we demonstrated how to compute clusters using the R program. Additionally, we described different methods for choosing the optimal number of clusters in a data set. These methods include the elbow, the silhouette and the gap statistic methods

#### **PART IV. Data reduction. Classification, Supervised learning.**

In this section, we apply the logistic regression to classify subjects according to whether the audits were productive or productive audits. Productive audits are those that result in an adjustment being made to a client's financial statement. The dependent variable or binary target variable for classification modeling (0/1), is TARGET\_Adjusted. The logistic procedure is limited to only two-class classification problems. There is an extension, called multinomial logistic regression, for multiclass classification problem.

Logistic regression is used to predict the class (or category) of individuals based on one or multiple predictor variables (x). It is used to model a binary outcome, that is a variable, which can have only two possible values: 0 or 1, yes or no. but in order to predict the probability that an audit will turn out productive or not, we must define a list of predictors or factors that may affect the probability of the dependent variable.

We chose the following features: “Age”, “Income”, “Hours”;

The model is as follow:

$$\text{Logit}(\text{TARGET\_Adjusted}) = B_0 + B_1 * \text{Age} + B_2 * \text{Income} + B_3 * \text{Hours} + \text{error};$$

Where  $\text{logit}(Y)$  denotes the log of the odds of  $Y=1$ ;

It should be noted that Logistic regression does not return directly the class of observations. It allows us to estimate the probability ( $p$ ) of class membership. The probability will range between 0 and 1. We choose the threshold probability at which the category flips from one to the other. By default, this is set to  $p = 0.5$ .

### *1. Supervised Learning*

Classification is a form of supervised learning where we have prior information on what the result (or class/group) should look like. In our case, we know the class label of the examples of the training set (0/1). Building a logistic ‘classifier’ is nothing more than running the logistic regression of TARGET\_Adjusted on Age, Income and Hours. In so doing we hope to determine which features are important for the classification, and which values of those features lead to success or failure. Once such values are determined based on the known cases, we have a rule or “classifier” that can be used to predict the outcome/result for new cases.

### *2. Methodology*

The data has already been loaded from the previous section. Moreover, all required R packages have been included before the analysis. We first partitioned our data set into two: testing and training. We also installed a few extra packages.

```
> install.packages("caret",
+                   repos = "http://cran.r-project.org",
+                   dependencies = c("Depends", "Imports", "Suggests"))

install.packages("gower");
library(gower);
library(caret);
library("mgcv");
library(readr);

> inTrain=createDataPartition (Dep,p=.75,list = FALSE);
> df=AuditData;
> training=df[inTrain,];
> testing=df[-inTrain,];
> nrow(training)
[1] 1500
> nrow(testing)
[1] 500
```

### *3. Running/Building the logistic classifier (Training)*

On the training set, we built the classifier using the generalized linear model procedure:

```
> Dep1=training$TARGET_Adjusted;
> set.seed(123);
> model1=glm(Dep1~training$Age+training$Income+training$Hours, family = binomial);
> model1

Call: glm(formula = Dep1 ~ training$Age + training$Income + training$Hours,
family = binomial)

Coefficients:
(Intercept)     training$Age   training$Income   training$Hours
-4.075e+00      3.745e-02     -6.844e-06      4.253e-02

Degrees of Freedom: 1499 Total (i.e. Null); 1496 Residual
Null Deviance: 1591
Residual Deviance: 1408      AIC: 1416
```

The estimated mode is as follow:

Logit (TARGET\_Adjusted) =  $-4.07 + 0.03745 \cdot \text{Age} - 0.000 \cdot \text{Income} + 0.0425 \cdot \text{Hours}$

As shown below, all predictors are significant at the 5% level (p-value < 0.05).

```
> summary(model1);

Call:
glm(formula = Dep1 ~ training$Age + training$Income + training$Hours,
family = binomial)

Deviance Residuals:
    Min      1Q      Median      3Q      Max
-2.0137 -0.7337 -0.5226 -0.2294  3.0522

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -4.075e+00 3.969e-01 -10.266 < 2e-16 ***
training$Age  3.745e-02 5.250e-03  7.132 9.88e-13 ***
training$Income -6.844e-06 1.332e-06 -5.137 2.79e-07 ***
training$Hours  4.253e-02 5.908e-03  7.198 6.10e-13 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1590.8 on 1499 degrees of freedom
Residual deviance: 1408.2 on 1496 degrees of freedom
AIC: 1416.2

Number of Fisher Scoring iterations: 5

> model1$coefficients
(Intercept)     training$Age   training$Income   training$Hours
-4.074940e+00   3.744654e-02    -6.844244e-06   4.252753e-02
> model1$df.null
[1] 1499
```

```

> model1$df.residual
[1] 1496
> model1>null.deviance
[1] 1590.787
> model1$deviance
[1] 1408.202

```

#### 4. Making predictions

The classifier can be used to make predictions we can test the classifier quality on the testing set.

```

> # Make predictions
> probabs=model1%>%predict(testing, type = "response");
> m.prob=predict(model1, data=testing, type="response");
> predicted.class=ifelse(probabs>0.5, "pos", "neg");

```

The predicted probabilities are (for a few sample of the testing set):

```

> probabs
      1       2       3       4       5       6
7 0.462543181 0.121111309 0.440109690 0.455407546 0.436613330 0.186452845 0.090
616226 0.115404115
      8       9      10      11      12      13      14
15 0.058218761 0.427423500 0.058383197 0.082219334 0.159664203 0.211450328 0.220
812616 0.458738311
      16      17      18      19      20      21      22
23 0.288381909 0.092887833 0.266937599 0.222050305 0.129543414 0.417079423 0.344
346201 0.150647028
      24      25      26      27      28      29      30
31 0.174426341 0.130899159 0.052173389 0.230361611 0.244928417 0.151676983 0.057
487107 0.294920368
      32      33      34      35      36      37      38
39 0.173836039 0.364188762 0.263951031 0.175819308 0.480775551 0.257605146 0.298
529020 0.284986372

```

The predicted classes are:

```

> predicted.class
      1       2       3       4       5       6       7       8       9       10      11      12      13
14 "neg" "neg"
      15      16      17      18      19      20      21      22      23      24      25      26      27
"neg" "neg"
      28      29      30      31      32      33      34      35      36      37      38      39
"neg" "neg"
      40      41      42      43      44      45      46      47      48      49      50      51
"neg" "neg"
      52      53      54      55      56      57      58      59      60      61      62      63      64
65 "neg" "neg"
      66      67      68      69      70      71      72      73      74      75      76      77      78
"neg" "neg"
      79      80      81      82      83      84      85      86      87      88      89      90      91
"neg" "neg"
      92      93      94      95      96      97      98      99      100      101      102      103      104
"neg" "neg"
      105      106      107      108      109      109      110      111      112      113      114      115      116
"neg" "neg"
      117      118      119      120      121      122      123      124      125      126      127      128      129
"neg" "neg"
      130      131      132      133      134      135      136      137      138      139      140      141      142
"neg" "neg"
      143      144      145      146      147      148      149      150      151      152      153      154      155
"neg" "neg"
      156      157      158      159      160      161      162      163      164      165      166      167      168
"neg" "neg"
      169      170      171      172      173      174      175      176      177      178      179      180      181
"neg" "neg"
      182      183      184      185      186      187      188      189      190      191      192      193      194
"neg" "neg"
      195      196      197      198      199      200      201      202      203      204      205      206      207
"neg" "neg"
      208      209      210      211      212      213      214      215      216      217      218      219      220
"neg" "neg"
      221      222      223      224      225      226      227      228      229      230      231      232      233
"neg" "neg"
      234      235      236      237      238      239      240      241      242      243      244      245      246
"neg" "neg"
      247      248      249      250      251      252      253      254      255      256      257      258      259
"neg" "neg"
      260      261      262      263      264      265      266      267      268      269      270      271      272
"neg" "neg"
      273      274      275      276      277      278      279      280      281      282      283      284      285
"neg" "neg"
      286      287      288      289      290      291      292      293      294      295      296      297      298
"neg" "neg"
      299      300      301      302      303      304      305      306      307      308      309      310      311
"neg" "neg"
      312      313      314      315      316      317      318      319      320      321      322      323      324
"neg" "neg"
      325      326      327      328      329      330      331      332      333      334      335      336      337
"neg" "neg"
      338      339      340      341      342      343      344      345      346      347      348      349      350
"neg" "neg"
      351      352      353      354      355      356      357      358      359      360      361      362      363
"neg" "neg"
      364      365      366      367      368      369      370      371      372      373      374      375      376
"neg" "neg"
      377      378      379      380      381      382      383      384      385      386      387      388      389
"neg" "neg"
      390      391      392      393      394      395      396      397      398      399      400      401      402
"neg" "neg"
      403      404      405      406      407      408      409      410      411      412      413      414      415
"neg" "neg"
      416      417      418      419      420      421      422      423      424      425      426      427      428
"neg" "neg"
      429      430      431      432      433      434      435      436      437      438      439      440      441
"neg" "neg"
      442      443      444      445      446      447      448      449      450      451      452      453      454
"neg" "neg"
      455      456      457      458      459      460      461      462      463      464      465      466      467
"neg" "neg"
      468      469      470      471      472      473      474      475      476      477      478      479      480
"neg" "neg"
      481      482      483      484      485      486      487      488      489      490      491      492      493
"neg" "neg"
      494      495      496      497      498      499      500      501      502      503      504      505      506
"neg" "neg"
      507      508      509      510      511      512      513      514      515      516      517      518      519
"neg" "neg"
      520      521      522      523      524      525      526      527      528      529      530      531      532
"neg" "neg"
      533      534      535      536      537      538      539      540      541      542      543      544      545
"neg" "neg"
      546      547      548      549      550      551      552      553      554      555      556      557      558
"neg" "neg"
      559      560      561      562      563      564      565      566      567      568      569      570      571
"neg" "neg"
      572      573      574      575      576      577      578      579      580      581      582      583      584
"neg" "neg"
      585      586      587      588      589      590      591      592      593      594      595      596      597
"neg" "neg"
      598      599      600      601      602      603      604      605      606      607      608      609      610
"neg" "neg"
      611      612      613      614      615      616      617      618      619      620      621      622      623
"neg" "neg"
      624      625      626      627      628      629      630      631      632      633      634      635      636
"neg" "neg"
      637      638      639      640      641      642      643      644      645      646      647      648      649
"neg" "neg"
      650      651      652      653      654      655      656      657      658      659      660      661      662
"neg" "neg"
      663      664      665      666      667      668      669      670      671      672      673      674      675
"neg" "neg"
      676      677      678      679      680      681      682      683      684      685      686      687      688
"neg" "neg"
      689      690      691      692      693      694      695      696      697      698      699      700      701
"neg" "neg"
      702      703      704      705      706      707      708      709      710      711      712      713      714
"neg" "neg"
      715      716      717      718      719      720      721      722      723      724      725      726      727
"neg" "neg"
      728      729      730      731      732      733      734      735      736      737      738      739      740
"neg" "neg"
      741      742      743      744      745      746      747      748      749      750      751      752      753
"neg" "neg"
      754      755      756      757      758      759      760      761      762      763      764      765      766
"neg" "neg"
      767      768      769      770      771      772      773      774      775      776      777      778      779
"neg" "neg"
      780      781      782      783      784      785      786      787      788      789      790      791      792
"neg" "neg"
      793      794      795      796      797      798      799      800      801      802      803      804      805
"neg" "neg"
      806      807      808      809      810      811      812      813      814      815      816      817      818
"neg" "neg"
      819      820      821      822      823      824      825      826      827      828      829      830      831
"neg" "neg"
      832      833      834      835      836      837      838      839      840      841      842      843      844
"neg" "neg"
      845      846      847      848      849      850      851      852      853      854      855      856      857
"neg" "neg"
      858      859      860      861      862      863      864      865      866      867      868      869      869

```

```

69   70   71   72   73   74   75   76   77   78   79   80   81
82 69   83   84   85   72   73   74   75   76   77   78   79   80   81
"neg" "neg" "neg" "neg" "neg" "neg" "pos" "neg" "neg" "pos" "neg" "neg" "neg"
"neg" "neg" "neg" "neg"

```

### 5. Comparing prediction and Observation

We can set a confusion table (contingency table 2x2) to compare the error rate on the testing set.

```

> m.pred=rep("0", dim(training)[1]);
> m.pred[m.prob>.5]="1";
> m.pred

[1] "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
"0" "0" "0" "0" "0" "0" "0" "[25]" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
"0" "0" "0" "0" "0" "0" "0" "[49]" "0" "1" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
"0" "0" "0" "0" "0" "0" "0" "[73]" "0" "0" "1" "0" "0" "1" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
"0" "0" "0" "0" "0" "0" "0" "[97]" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
"0" "0" "0" "0" "0" "0" "0" "[121]" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "1"
"0" "0" "0" "0" "0" "0" "0" "[145]" "1" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
"0" "0" "0" "0" "0" "0" "0" "[169]" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
"0" "0" "0" "0" "0" "0" "0" "[193]" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
"0" "0" "0" "0" "0" "0" "0" "[217]" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
"0" "0" "0" "0" "0" "0" "0"

> table(m.pred, Dep1);
    Dep1
m.pred 0     1
      0 1139  303
      1   27   31

```

The total prediction rate of success is  $(1139+31)/(27+303+1139+31) = 0.78 = 78\%$ . Our classifier has about 80% success rate on the testing set and 20% error rate.

### Conclusion:

We used a logistic regression model to built binary classifier for the variable TARGET\_Adjusted. The model predicts whether the audits were productive or productive audits using the estimated equation: Logit (TARGET\_Adjusted) = - 4.07 + 0.03745\*Age - 0.000\*Income + 0.0425\*Hours. This model has an 80% success rate on new data. It can be used to predict future outcome for new data.

We used a logistic regression model to build a binary classifier for the variable TARGET\_Adjusted. The model predicts whether the audits were productive or productive audits using the estimated equation: Logit(TARGET\_Adjusted) = 4.07 + 0.03745\*Age - 0.000\*Income + 0.0425\*Hours. This model has an 80% success rate on new data. It can be used to predict future outcome for new data.

## References

1. Hartigan, JA, and MA Wong. 1979. “Algorithm AS 136: A K-means clustering algorithm.” *Applied Statistics*. Royal Statistical Society, 100–108.
2. MacQueen, J. 1967. “Some Methods for Classification and Analysis of Multivariate Observations.” In *Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability, Volume 1: Statistics*, 281–97. Berkeley, Calif.: University of California Press. <http://projecteuclid.org:443/euclid.bsmsp/1200512992>.
3. Charrad, Malika, Nadia Ghazzali, Véronique Boiteau, and Azam Niknafs. 2014. “NbClust: An R Package for Determining the Relevant Number of Clusters in a Data Set.” *Journal of Statistical Software* 61: 1–36. <http://www.jstatsoft.org/v61/i06/paper>.